

# Variational Oblique Predictive Clustering Trees

Viktor Andonovikj

Jožef Stefan Institute & Jožef Stefan International Postgraduate School

September 26, 2024

# Outline

- Key Idea
- Introduction to Structured Output Prediction
- Variational SPYCT
- Experimental Setting
- Data
- Results
- Interpretability
- Conclusion

# Key Idea

- We focus on **decision tree** type of models - popular due to their interpretability and simplicity.
- We aim to combine the **predictive power** of **ensemble** methods with the **interpretability** of a **single** decision tree.
- **Variational SPYCT** incorporates Bayesian inference to enhance both predictive performance and decision-making transparency.

# Introduction to Structured Output Prediction (SOP)

- **Structured Output Prediction (SOP)** involves predicting multiple interdependent outputs.
- SOP tasks include simultaneous prediction of:
  - ▶ Multiple continuous values.
  - ▶ Multiple discrete values.
  - ▶ Hierarchically organized discrete values.
- **Real-world applications:**
  - ▶ Drug discovery: predicting multiple biological properties of molecules.
  - ▶ Recommender systems: predicting user preferences across multiple items.
  - ▶ Genomics: predicting gene functions based on interdependent traits.

# Predictive Clustering Framework

- **Predictive Clustering:** Combines clustering and prediction by treating prediction as a hierarchical clustering task where similar instances are grouped and predictions are made for each cluster.
- **Characteristics:**
  - ▶ Supports both supervised and semi-supervised learning.
  - ▶ State-of-the-art performance via ensemble learning.
  - ▶ Offers interpretable models through feature importance analysis.
  - ▶ Provides a unified framework for predictive modeling across multiple tasks.

# Predictive Clustering Trees (PCT) vs Oblique Predictive Clustering Trees (SPYCT)




Figure: A) Learned split - SPYCT in blue, PCT in red B) PCT C) SPYCT [1]

- **PCT**: Uses axis-aligned splits (based on single features), fully interpretable models.
- **SPYCT**: Uses oblique splits (linear combinations of features), more flexibility for high-dimensional and sparse data, suited for more complex tasks with intricate decision boundaries.

# Introduction to Variational SPYCT

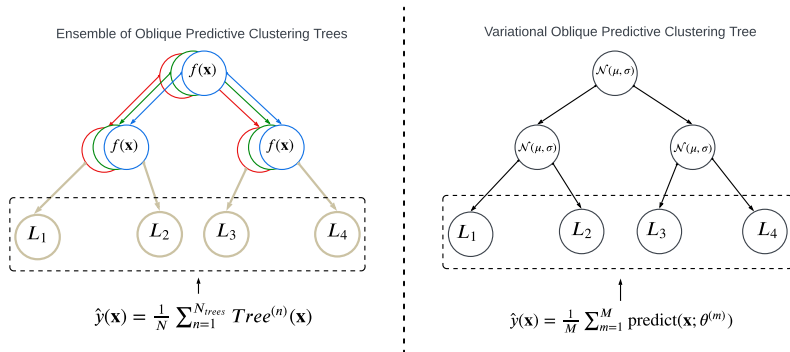


Figure: Ensemble of SPYCTs vs Variational SPYCT

# Introduction to Variational SPYCT

- **Motivation:** Variational SPYCT (VSPYCT) integrates variational Bayes for improved decision-making within a single model, eliminating the need for ensembles.
- **Uncertainty quantification:** Embeds Bayesian inference directly into the decision tree structure, providing insight into decision processes and confidence levels.
- **Novelty:** VSPYCT introduces probabilistic treatment of oblique splits, offering a paradigm shift toward interpretable, and reliable machine learning models.



# Optimization through Variational Bayes

- **Bayes' Theorem:**

$$p(\theta|x) = \frac{p(x|\theta)p(\theta)}{p(x)}$$

- The primary computational challenge lies in calculating the posterior  $p(\theta|x)$ , where the denominator  $p(x)$  requires:

$$p(x) = \int_{\theta} p(x|\theta)p(\theta) d\theta$$

- **Variational Bayes (VB)** approximates the posterior with  $q_{\omega^*}(\theta)$ , minimizing the KL divergence:


$$\omega^* = \arg \min_{\omega \in \Omega} KL(q_{\omega}(\theta) \parallel p(\theta|x))$$

- The Evidence Lower Bound (ELBO) is maximized to improve the approximation:

$$KL(q_{\omega}(\theta) \parallel p(\theta|x)) = -\mathbb{E}_q[\log p(x, \theta) - \log q_{\omega}(\theta)] + \log p(x)$$

- VB's computational efficiency surpasses methods like MCMC, though it introduces bias based on the choice of the variational family  $\mathcal{Q}$ .

# Optimization through Variational Bayes



**Figure:** Optimisation process of finding the closest variational distribution  $q_\omega(\theta)$  over the set of latent variables  $\omega$ .

# Methodology Overview

- VSPYCT follows SPYCT's tree-based architecture, but replaces fixed parameters with random variables.
- **Key difference:** VSPYCT uses variational Bayes (VB) for probabilistic split optimization, improving handling of noisy data and uncertainty.
- Split parameters (weights  $\mathbf{w}$  and bias  $b$ ) are modeled as random variables, allowing uncertainty estimation.

$$f(\mathbf{x}) = \sigma \left( \mathbf{w}^\top \mathbf{x} + b \right)$$

# Learning Splits through Variational Bayes

- Weights  $\mathbf{w}$  and bias  $b$  have Gaussian priors:

$$\mathbf{w} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}), \quad b \sim \mathcal{N}(0, 1)$$

- Variational Bayes approximates the posterior distribution with:

$$q(\mathbf{w}, b | \mathcal{D}) = \mathcal{N}(\mathbf{w} | \boldsymbol{\mu}_w, \boldsymbol{\Sigma}_w) \mathcal{N}(b | \mu_b, \sigma_b^2)$$

- ELBO maximization:

$$\mathcal{L}(\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \mathbb{E}_{q(\mathbf{w}, b | \mathcal{D})}[\log p(\mathcal{D} | \mathbf{w}, b)] - KL(q(\mathbf{w}, b | \mathcal{D}) \parallel p(\mathbf{w}, b))$$

# Algorithm for Learning a Split

---

**Algorithm 1** Variational Learning of Split Parameters

---

```
1: Input:  $\mathcal{D} = \{\mathbf{X}, \mathbf{Y}\}$ ,  $\theta = \{\mathbf{w}_0, b_0\}$ ,  $E$  (epochs),  $\lambda$  (learning rate),  $\beta$  (batch size),  $\sigma$  (selection probability)
2: Output:  $\Theta = \{\mu_w, \Sigma_w, \mu_b, \sigma_b^2\}$  (variational parameters)
3: procedure LEARN_SPLIT( $\mathcal{D}, \theta, E, \lambda, \beta, \sigma$ )
4:   Initialize  $\Theta = \{\mu_w, \Sigma_w, \mu_b, \sigma_b^2\}$ 
5:   for  $e \in \{1, \dots, E\}$  do
6:     for each mini-batch  $\mathcal{B} \subseteq \mathcal{D}$  of size  $\beta$  do
7:       Sample  $\mathbf{w} \sim \mathcal{N}(\mu_w, \Sigma_w)$ ,  $b \sim \mathcal{N}(\mu_b, \sigma_b^2)$ 
8:       Compute impurity  $\Omega(\mathbf{X}, \mathbf{Y}; \mathbf{w}, b)$ 
9:       Compute ELBO  $\mathcal{L}(\mathcal{B}; \Theta)$ 
10:      Update  $\Theta$ 
11:    end for
12:  end for
13:  return  $\Theta$ 
14: end procedure
```

# Deriving the Impurity Function

- For a given split node  $i$ , we define the fuzzy membership:

$$FM_i = \sigma(\mathbf{x}^\top \mathbf{w}_i + b_i), \quad 1 - FM_i = \text{right group membership}$$

- The split fitness function is the following:

$$f(\mathbf{w}_i, b_i) = Z \cdot \text{imp}(FM) + (L + U - Z) \cdot \text{imp}(1 - FM)$$

- Impurity is given by the variances of the target  $Y$  and features  $X$ :

$$\text{imp}(FM) = \sum_{k=1}^N \sigma_{\tilde{X}_k}^2 + \sum_{k=1}^T \sigma_{\tilde{Y}_k}^2$$

- In VSPYCT, we minimize this impurity by observing a target impurity of  $\text{impurity}/2$  at each step using Variational Bayes.

# Making a Prediction

- Prediction involves traversing the tree, making probabilistic decisions at each node.
- The split is determined by evaluating the function:

$$f(\mathbf{x}) = \sigma \left( \mathbf{w}^\top \mathbf{x} + b \right)$$

- The final prediction  $\hat{y}$  is the prototype value at the reached leaf:

$$\hat{y} = \frac{1}{|\mathbf{Y}|} \sum_{i=1}^{|\mathbf{Y}|} y_i$$

# Prediction Process in VSPYCT

---

**Algorithm 2** Prediction using Monte Carlo Sampling

---

- 1: **Input:** Feature vector  $\mathbf{x}$ , tree  $T$ , samples  $M$
  - 2: **Output:** Prediction  $\hat{y}$
  - 3: Initialize  $\hat{y}_{\text{sum}} = 0$
  - 4: **for**  $m = 1$  to  $M$  **do**
  - 5:     Traverse  $T$  from root to leaf with sampled  $\mathbf{w}^{(m)}, b^{(m)}$
  - 6:      $\hat{y}_{\text{sum}} \leftarrow \hat{y}_{\text{sum}} + \text{prediction at leaf}$
  - 7: **end for**
  - 8: **return**  $\hat{y} = \frac{\hat{y}_{\text{sum}}}{M}$
-



# Feature Importance

- Feature importance in VSPYCT: Calculated as the influence of features across all splits.
- The importance of a feature is determined by:

$$\text{Imp}(T) = \sum_{s \in T} \left( \frac{s_n}{N} \right) \left( \frac{\mathbb{E}[\mathbf{w}_s]}{|\mathbb{E}[\mathbf{w}_s]|_1} \right)$$

- Weights  $\mathbf{w}$  are sampled from the posterior distribution, and importance is aggregated over all splits.

# Time Complexity Analysis

- Time complexity of learning a split in VSPYCT:

$$\mathcal{O}(MNI_{vb}(D + K))$$

- $M$ : Number of Monte Carlo samples,  $N$ : Data points,  $D$ : Features,  $K$ : Clustering attributes,  $I_{VB}$ : Number of optimization iterations.

# Experimental Setting

- **Comparison:** VSPYCT vs SPYCT (single tree and ensemble).
- **Tasks:** Single-target and multi-target regression, binary and multi-class classification.
- **Goal:** Evaluate predictive performance of VSPYCT across different modeling scenarios.

# Data

- Datasets cover regression and classification tasks.
- Number of examples (N), features (D), targets (T), and classes (C).

**Table:** Summary of the datasets used in the experiments.

Dataset	Type of Task	N	D	T	C
rf1 [2]	Multi-target regression	9125	64	8	–
rf2 [2]	Multi-target regression	9125	576	8	–
atp1d [2]	Multi-target regression	337	411	6	–
atp7d [2]	Multi-target regression	296	411	6	–
scm1d [2]	Multi-target regression	9803	280	16	–
house_8L [3]	Single-target regression	22784	8	1	–
puma8NH [3]	Single-target regression	8192	8	1	–
diabetes [3]	Binary classification	768	8	1	2
banknote [3]	Binary classification	1372	4	1	2
gas-drift [3]	Multi-class classification	13910	128	1	6
balance [3]	Multi-class classification	625	4	1	3

# Evaluation Metrics

- **Regression:** Mean Absolute Error (MAE).

$$MAE = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$$

- **Classification:** F1 score (macro-averaged for multi-class tasks).
- **Cross-validation:** 5-fold cross-validation used to reduce bias.

# Regression Results

Table: MAE Scores for Regression Datasets

Dataset	SPYCT-single tree	SPYCT-ensemble	VSPYCT
rf1	29.38	<b>29.34</b>	29.37
rf2	29.37	29.49	<b>29.36</b>
atp1d	95.17	<b>70.93</b>	77.24
atp7d	115.54	<b>73.08</b>	91.28
scm1d	205.99	206.01	<b>205.72</b>
house_8L	25378.02	<b>19817.56</b>	22992.24
puma8NH	2.80	<b>2.62</b>	2.74

- VSPYCT outperforms or closely matches the SPYCT ensemble on several datasets (rf2, scm1d).
- Ensemble methods generally achieve lower MAE, but VSPYCT provides competitive results.

# Classification Results

Table: F1 Scores for Classification Datasets

Dataset	SPYCT-single tree	SPYCT-ensemble	VSPYCT
diabetes	0.53	<b>0.60</b>	0.59
banknote	0.97	<b>0.99</b>	0.98
gas-drift	0.98	0.99	<b>0.99</b>
balance	0.60	0.66	<b>0.73</b>


- VSPYCT achieves competitive performance across datasets, often matching or surpassing ensembles.
- In the balance dataset, VSPYCT achieves the best F1 score.

# Interpretability - An example with real dataset

- **Dataset:** Unemployment dataset from the Slovenian Public Employment Service (PES), consisting of 74,086 anonymized instances.
- **Features:** Age, gender, education, work experience, and other personal/professional characteristics.
- **Target:** Time until the jobseeker becomes employed or exits the study (measured in days). Some records are right-censored (event not observed).
- **Goal:** Predict the time-to-event (employment) with censored data, using a semi-supervised learning approach for handling the inherent missing information.
- **Challenges:** Diverse attribute types (categorical, numerical, temporal) and the presence of censored data.




# Structure of the learned VSPYCT model



**Figure:** VSPYCT tree structure. Colors represent different feature groups (e.g., education, profession).


# Predicted Survival Curve



**Figure:** Predicted survival curve for a jobseeker. Blue: mean prediction; shaded area: confidence interval; vertical line: true event time.

- The confidence interval captures the uncertainty behind the predictions.

# Feature Importance



**Figure:** Feature importance from the VSPYCT model. "Months of work experience" is the most important factor in the model's decisions.

# Conclusion

- **VSPYCT Model:** Combines oblique splits, variational Bayes, and Bayesian inference in an interpretable predictive framework.
- **Key Characteristic:** Balances the performance of ensemble models with the interpretability of single decision trees, while introducing uncertainty quantification.
- **Results:** Competitive performance, sometimes surpassing ensembles of SPYCT.
- **Uncertainty Quantification:** Enhances decision-making reliability, making the model applicable to domains where certainty in predictions is critical.
- **Future Work:** Further performance improvements and extending applications to diverse predictive tasks.

# References I

- [1] T. Stepišnik and D. Kocev, "Oblique predictive clustering trees," *Knowledge-Based Systems*, vol. 227, p. 107 228, Sep. 2021, ISSN: 0950-7051. DOI: [10.1016/j.knosys.2021.107228](https://doi.org/10.1016/j.knosys.2021.107228). [Online]. Available: <http://dx.doi.org/10.1016/j.knosys.2021.107228>.
- [2] Mulan, *Mulan: A java library for multi-label learning*, <http://mulan.sourceforge.net/datasets.html>, Accessed: 2020-04-15.
- [3] OpenML, *Openml*, <https://www.openml.org>, Accessed: 2020-04-15.

# Thank You

Questions?